

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
from matplotlib import pyplot as plt
from sklearn.preprocessing import scale
import warnings
warnings.filterwarnings('ignore')
```

## Load the Dataset into the tool

```
In [11]: df=pd.read_csv("abalone.csv")
df.head()
```

```
Out[11]:
```

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	15
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	7
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	9
3	M	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	10
4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	7

```
In [13]: df.shape
```

```
Out[13]: (4177, 9)
```

```
In [15]: Age=1.5+df.Rings
df["Age"]=Age
df=df.rename(columns = {'Whole weight':'Whole_weight','Shucked weight': 'Shucked_weight',
                        'Shell weight': 'Shell_weight'})
df=df.drop(columns=["Rings"],axis=1)
df.head()
```


```
Out[15]:
```

	Sex	Length	Diameter	Height	Whole_weight	Shucked_weight	Viscera_weight	Shell_weight	Age
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	16.5
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	8.5
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	10.5
3	M	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	11.5
4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	8.5

## Univariate Analysis

```
In [16]: sns.displot(df["Age"], color='darkorange')
```


```
Out[16]: <seaborn.axisgrid.FacetGrid at 0x2629a2418d0>
```



## Bi - Variate Analysis


```
In [17]: sns.scatterplot(x=df.Age,y=df.Height,color='black')
```

```
Out[17]: <AxesSubplot: xlabel='Age', ylabel='Height'>
```



```
In [8]: sns.pairplot(df)
```

```
Out[8]: <seaborn.axisgrid.PairGrid at 0x2629577ce20>
```



Perform descriptive statistics on the dataset

```
In [18]: df.describe()
```

Out[18]:		Length	Diameter	Height	Whole_weight	Shucked_weight	Viscera_weight	Shell_weight	
	count	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4
	mean	0.523992	0.407881	0.139516	0.828742	0.359367	0.180594	0.238831	
	std	0.120093	0.099240	0.041827	0.490389	0.221963	0.109614	0.139203	
	min	0.075000	0.055000	0.000000	0.002000	0.001000	0.000500	0.001500	
	25%	0.450000	0.350000	0.115000	0.441500	0.186000	0.093500	0.130000	
	50%	0.545000	0.425000	0.140000	0.799500	0.336000	0.171000	0.234000	
	75%	0.615000	0.480000	0.165000	1.153000	0.502000	0.253000	0.329000	
	max	0.815000	0.650000	1.130000	2.825500	1.488000	0.760000	1.005000	

Check for Missing values and deal with them.


```
In [19]: df.isna().sum()
```

```
Out[19]: Sex                0
Length                0
Diameter              0
Height                0
Whole_weight          0
Shucked_weight        0
Viscera_weight        0
Shell_weight          0
Age                   0
dtype: int64
```

Find the outliers and replace them outliers.

```
In [20]: sns.boxplot(df['Length'])
```

```
Out[20]: <AxesSubplot: >
```



Check for Categorical columns and perform encoding.

```
In [21]: df['Sex'] = df['Sex'].replace({'M':1, 'F':0, 'I':2}, inplace=True)
```

```
df.tail()
```

```
Out[21]:
```

	Sex	Length	Diameter	Height	Whole_weight	Shucked_weight	Viscera_weight	Shell_weight	Age
4172	0	0.565	0.450	0.165	0.8870	0.3700	0.2390	0.2490	12.5
4173	1	0.590	0.440	0.135	0.9660	0.4390	0.2145	0.2605	11.5
4174	1	0.600	0.475	0.205	1.1760	0.5255	0.2875	0.3080	10.5
4175	0	0.625	0.485	0.150	1.0945	0.5310	0.2610	0.2960	11.5
4176	1	0.710	0.555	0.195	1.9485	0.9455	0.3765	0.4950	13.5

## Split the data into dependent and independent variables.

```
In [24]: outliers=df.quantile(q=(0.25,0.75))
outliers
```

```
Out[24]:
```

	Sex	Length	Diameter	Height	Whole_weight	Shucked_weight	Viscera_weight	Shell_weight	Age
0.25	0.0	0.450	0.35	0.115	0.4415	0.186	0.0935	0.130	9.5
0.75	2.0	0.615	0.48	0.165	1.1530	0.502	0.2530	0.329	12.5

```
In [25]: a = df.Age.quantile(0.25)
b = df.Age.quantile(0.75)
c = b - a
lower_limit = a - 1.5 * c
df.median(numeric_only=True)
```

```
Out[25]: Sex                1.0000
Length              0.5450
Diameter            0.4250
Height              0.1400
Whole_weight        0.7995
Shucked_weight      0.3360
Viscera_weight      0.1710
Shell_weight        0.2340
Age                 10.5000
dtype: float64
```

```
In [27]: from sklearn.preprocessing import LabelEncoder

lab = LabelEncoder()
df.Sex = lab.fit_transform(df.Sex)

df.head()
```

```
Out[27]:
```

	Sex	Length	Diameter	Height	Whole_weight	Shucked_weight	Viscera_weight	Shell_weight	Age
0	1	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	16.5
1	1	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	8.5
2	0	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	10.5
3	1	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	11.5
4	2	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	8.5

# Scale the independent variables

```
In [28]: x=df.drop(columns=['Length', 'Height'])
print(x)
```

	Sex	Diameter	Whole_weight	Shucked_weight	Viscera_weight	\
0	1	0.365	0.5140	0.2245	0.1010	
1	1	0.265	0.2255	0.0995	0.0485	
2	0	0.420	0.6770	0.2565	0.1415	
3	1	0.365	0.5160	0.2155	0.1140	
4	2	0.255	0.2050	0.0895	0.0395	
...	...	...	...	...	...	
4172	0	0.450	0.8870	0.3700	0.2390	
4173	1	0.440	0.9660	0.4390	0.2145	
4174	1	0.475	1.1760	0.5255	0.2875	
4175	0	0.485	1.0945	0.5310	0.2610	
4176	1	0.555	1.9485	0.9455	0.3765	

	Shell_weight	Age
0	0.1500	16.5
1	0.0700	8.5
2	0.2100	10.5
3	0.1550	11.5
4	0.0550	8.5
...	...	...
4172	0.2490	12.5
4173	0.2605	11.5
4174	0.3080	10.5
4175	0.2960	11.5
4176	0.4950	13.5

[4177 rows x 7 columns]

```
In [29]: x=scale(x)
print(x)
```

```
[[-0.0105225 -0.43214879 -0.64189823 ... -0.72621157 -0.63821689
  1.57154357]
 [-0.0105225 -1.439929   -1.23027711 ... -1.20522124 -1.21298732
 -0.91001299]
 [-1.26630752  0.12213032 -0.30946926 ... -0.35668983 -0.20713907
 -0.28962385]
 ...
 [-0.0105225  0.67640943  0.70821206 ...  0.97541324  0.49695471
 -0.28962385]
 [-1.26630752  0.77718745  0.54199757 ...  0.73362741  0.41073914
  0.02057072]
 [-0.0105225  1.48263359  2.28368063 ...  1.78744868  1.84048058
  0.64095986]]
```

## Split the data into training and testing.

```
In [42]: y = df["Sex"]
y.head()
```

```
Out[42]: 0    1
         1    1
         2    0
         3    1
         4    2
         Name: Sex, dtype: int64
```

```
In [36]: x=df.drop(columns=["Sex"], axis=1)
x.head()
```

```
Out[36]:
```

	Length	Diameter	Height	Whole_weight	Shucked_weight	Viscera_weight	Shell_weight	Age
0	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	16.5
1	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	8.5
2	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	10.5
3	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	11.5
4	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	8.5

```
In [43]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_state=0)
x_train.shape
```

```
Out[43]: (3341, 8)
```

```
In [44]: x_test.shape
```

```
Out[44]: (836, 8)
```

```
In [45]: y_test.shape
```

```
Out[45]: (836,)
```

## Build the Model.

```
In [46]: from sklearn.tree import DecisionTreeClassifier
model=DecisionTreeClassifier()
df = pd.get_dummies(df, drop_first=True)
df.head()
```

```
Out[46]:
```

	Sex	Length	Diameter	Height	Whole_weight	Shucked_weight	Viscera_weight	Shell_weight	Age
0	1	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	16.5
1	1	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	8.5
2	0	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	10.5
3	1	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	11.5
4	2	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	8.5

```
In [47]: X = df.drop('Height', axis=1)
y = df['Height']

from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33)

from sklearn.preprocessing import StandardScaler
ss = StandardScaler()

X_trains = ss.fit_transform(X_train)
X_tests = ss.transform(X_test)
```

```
In [48]: from sklearn.linear_model import LinearRegression
lr = LinearRegression()

lr.fit(X_trains, y_train)
pred = lr.predict(X_tests)

from sklearn.metrics import r2_score, roc_auc_score, mean_squared_error
rmse = np.sqrt(mean_squared_error(y_test, pred))
r2 = r2_score(y_test, pred)

print("The root mean Sq error calculated from the base model is:",rmse)
print("The r2-score is:",r2)
```

The root mean Sq error calculated from the base model is: 0.01738935058088849  
The r2-score is: 0.8102454999039992

## selecting best feautre

```
In [49]: from sklearn.feature_selection import RFE
lr = LinearRegression()
n = [{'n_features_to_select':list(range(1,10))}]
rfe = RFE(lr)

from sklearn.model_selection import GridSearchCV
gsearch = GridSearchCV(rfe, param_grid=n, cv=3)
gsearch.fit(X, y)

gsearch.best_params_
```

Out[49]: {'n\_features\_to\_select': 8}

```
In [50]: lr = LinearRegression()
rfe = RFE(lr, n_features_to_select=8)
rfe.fit(X,y)

pd.DataFrame(rfe.ranking_, index=X.columns, columns=['Class'])
```

Out[50]:

	Class
Sex	1
Length	1
Diameter	1
Whole_weight	1
Shucked_weight	1
Viscera_weight	1
Shell_weight	1
Age	1

## Measure the performance using Metrics

```
In [51]: from sklearn.neighbors import KNeighborsRegressor
from sklearn.ensemble import RandomForestRegressor
from sklearn.linear_model import LinearRegression
ensemble import GradientBoostingRegressor
```



```

from sklearn.linear_model import Ridge
from sklearn.svm import SVR
from sklearn import model_selection
from sklearn.model_selection import cross_val_predict

models = [    SVR(),
              RandomForestRegressor(),
              GradientBoostingRegressor(),
              KNeighborsRegressor(n_neighbors = 4)]

results = []
names = ['SVM', 'Random Forest', 'Gradient Boost', 'K-Nearest Neighbors']
for model, name in zip(models, names):
    kfold = model_selection.KFold(n_splits=10)
    cv_results = model_selection.cross_val_score(model, X_train, y_train, cv=kfold)
    rmse = np.sqrt(mean_squared_error(y, cross_val_predict(model, X , y, cv=3)))
    results.append(rmse)
    names.append(name)
    msg = "%s: %f" % (name, rmse)
    print(msg)

```

```

SVM: 0.036681
Random Forest: 0.025132
Gradient Boost: 0.023635
K-Nearest Neighbors: 0.023851

```

In [ ]: